


# STAT5003

## Week 2: Regression and Smoothing

Dr. Justin Wishart



# Readings and functions covered

- Introduction to Statistical Learning James, Witten, Hastie, and Tibshirani (2013)
  - Chapter 3 (Linear regression)
  - Chapter 7.4 to 7.6 (Smoothing)
-  functions
  - `outcome ~ feature1 + feature2` (formulae)
  - `lm` (Linear model)
  - `confint` (confidence intervals)
  - `subset` (argument and function)
  - `predict` (Make predictions from a model)

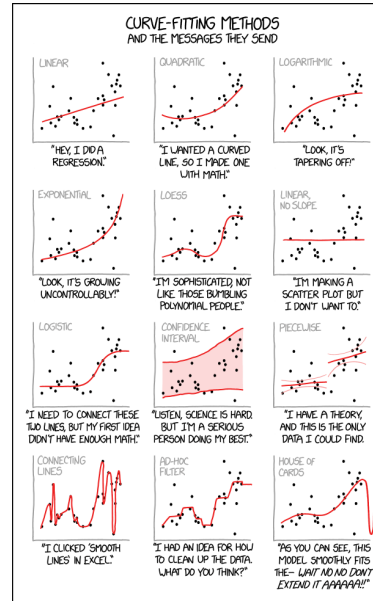
# Regression



THE UNIVERSITY OF  
**SYDNEY**

# Regression

- Numerically fitting the model is easy




- Knowing how to appropriately fit the model is where you add value.

Line of Best Fit

# The prediction problem

What is the price of a 100 sqm house in St Kilda?



# The linear regression model


$$Y = \beta_0 + \beta_1 X + \varepsilon$$

↓

$$y_i = b_0 + b_1 x_i + e_i$$

- $X$  is the **predictor** (feature or independent variable)
- $Y$  is the **response** (target or dependent variable)
- $\beta_0$  is the **intercept** of the regression line
  - Expected value of  $Y$  when  $X = 0$
- $\beta_1$  is the **slope** of the regression line
  - mean increase in  $Y$  for a *unit* increase in  $X$
- $\varepsilon$  is the **unexplained variation** or random error.
  - Classically assumed to be normally distributed with mean zero and finite variance.

# Performance of regression estimates




Candidate Line — Line of 'Best Fit' — True — Other possible lines

- Data was simulated from model  $Y = -0.5 + 0.75X + \varepsilon$
- True line shown in blue
- Standard linear regression fit shown in red
- Why not one of the green lines?



# How to determine the best estimates of $\beta_0 + \beta_1 X$ ?

- The notion of best needs a **criterion** to measure against.



- Easiest mathematical solution is the **least squares criterion**

- Minimise the residual sum of squares  $RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \sum_{i=1}^n e_i^2$

# Least squares equations

- Can show by simple calculus the following:

- Regression (slope) coefficient:  $b_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2} = \frac{\text{cov}(x, y)}{\text{var}(x)}$

- Intercept:  $b_0 = \bar{y} - b_1 \bar{x}$

- This leads to the estimated regression line:

$$\hat{y} = b_0 + b_1 x$$

- Least squares regression line since it minimises the residual sum of squares.

## Basic uses of Simple Linear Regression

# Prediction using `lm`

```
lm.fit <- lm(Price ~ BuildingArea, data = st.kilda.data)
summary(lm.fit)
```

```
##
## Call:
## lm(formula = Price ~ BuildingArea, data = st.kilda.data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -817415 -201614  -85181   19895 3403199
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -129484.0    91775.9  -1.411   0.161
## BuildingArea   11209.5     799.8  14.015 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 490300 on 99 degrees of freedom
## Multiple R-squared:  0.6649,    Adjusted R-squared:  0.6615
## F-statistic: 196.4 on 1 and 99 DF,  p-value: < 2.2e-16
```

# Standard error of population mean

- Consider single population estimation problem .
  - Wish to estimate some mean,  $\mu$ , of some random variable  $Y$  .
  - If  $Y_i$  is sampled then  $\hat{\mu} = \overline{Y}$  estimates  $\mu$  with

$$\text{Var}(\hat{\mu}) = (\text{SE}(\hat{\mu}))^2 = \frac{\sigma^2}{n}$$

- $\sigma^2$  is the variance of  $Y_i$
- $n$  is the sample size.

# Standard error of regression coefficient estimates

- Same concept applies to the regression estimates

$$SE(\hat{\beta}_0) = \sigma \sqrt{\frac{1}{n} + \frac{\bar{x}^2}{\sum_{i=1}^n (x_i - \bar{x})^2}}$$
$$SE(\hat{\beta}_1) = \frac{\sigma}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2}}$$

where  $\sigma^2 = \text{Var}(\varepsilon)$

- As  $n \rightarrow \infty$ ,  $SE(\hat{\beta}_0) \rightarrow 0$  and  $SE(\hat{\beta}_1) \rightarrow 0$
- Interestingly, if the  $x_i$  are more spread out, the standard errors will be smaller
  - more leverage to estimate the parameters.

# Using standard errors to compute confidence intervals


```
summary(lm.fit) # Truncated output with coefficient table
```

```
...  
## Coefficients:  
##           Estimate Std. Error t value Pr(>|t|)  
## (Intercept) -129484.0    91775.9  -1.411    0.161  
## BuildingArea  11209.5     799.8   14.015 <2e-16 ***  
## ---  
##  
## Residual standard error: 490300 on 99 degrees of freedom  
...
```

- We can use the standard error to estimate the 95% confidence interval as:
  - $(\hat{\beta}_1 - t_{n-2,0.975}SE(\hat{\beta}_1), \hat{\beta}_1 + t_{n-2,0.975}SE(\hat{\beta}_1)) = b_1 \pm t_{n-2,0.975}SE(b_1) = b_1 \pm t_{99,0.975}SE(b_1)$
- In our housing example, the 95% confidence interval for the coefficient of BuildingArea is [9622.6968, 12796.3032]

$$b_1 \pm t_{n-2,0.975}SE(b_1) = 1.12095 \times 10^4 \pm 1.984 \times 799.8 = (9622.6968, 12796.3032)$$

# Confidence intervals of regression coefficients

- More directly in  code.
  - Use the `confint` function.

```
confint(lm.fit)
```

```
##              2.5 %    97.5 %  
## (Intercept) -311587.233 52619.18  
## BuildingArea   9622.491 12796.50
```

- This is exact and no precision lost to rounding error.
- Easy to change confidence level (99% below)

```
confint(lm.fit, level = 0.99)
```

```
##              0.5 %    99.5 %  
## (Intercept) -370524.63 111556.57  
## BuildingArea   9108.86 13310.13
```



# Is BuildingArea a good predictor of price?

```
summary(lm.fit) # truncated for coefficient table
```

```
...  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)  
## (Intercept) -129484.0    91775.9  -1.411    0.161  
## BuildingArea  11209.5     799.8   14.015 <2e-16 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
...
```

- Linear regression assumes  $Y = \beta_0 + \beta_1 X + \varepsilon$
- If BuildingArea is not linearly related to Price, then  $\beta_1 = 0$ .
- Can conduct a test of significance  $H_0 : \beta_1 = 0$  against  $H_1 : \beta_1 \neq 0$
- Can conduct a hypothesis test by computing the t-statistic:

$$t = \frac{\hat{\beta}_1 - \beta_1}{SE(\hat{\beta}_1)} \stackrel{H_0}{=} \frac{\hat{\beta}_1}{SE(\hat{\beta}_1)}$$

# Is BuildingArea a good predictor of price?

```
summary(lm.fit) # truncated for coefficient table
```

```
...  
## Coefficients:  
##           Estimate Std. Error t value Pr(>|t|)  
## (Intercept) -129484.0    91775.9  -1.411    0.161  
## BuildingArea  11209.5     799.8   14.015 <2e-16 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
...
```

- The **p-values** for each significance test in the last column.
- Recall, **p-value** gives the probability of observing your test statistic (and other scenarios support  $H_1$ ) assuming  $H_0$  is true.
- Small p-value here gives very little evidence to support the claim that there is no relationship between Price and BuildingArea

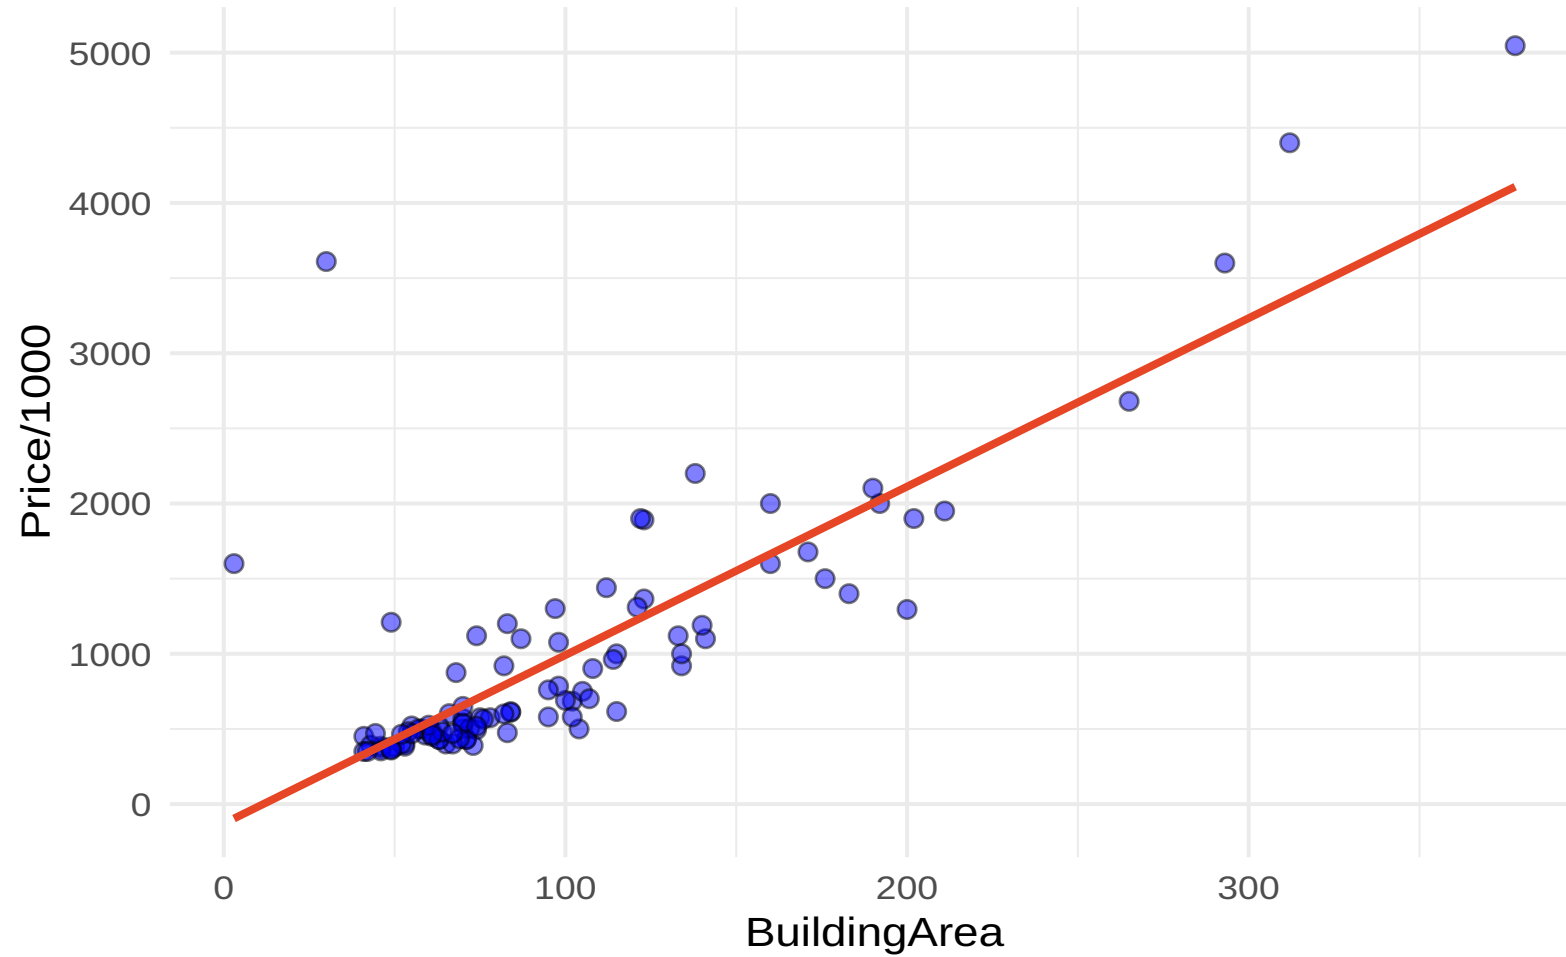
# Goodness of fit statistic

- Goodness of fit is measured by the **coefficient of determination** or  $R^2$

$$R^2 = \frac{\text{Total Sum of Squares} - \text{Residual Sum of Squares}}{\text{Total Sum of Squares}}$$
$$= \frac{\sum_{i=1}^n (y_i - \bar{y})^2 - \sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

- $R^2$  is a measure between 0 and 1
- It measures the proportion of variation in the response Y , explained by the linear regression on X
  - A value of 0 indicates **none** of the variance in Y can be explained linearly by X
  - A value of 1 indicates **all** of the variance in Y can be explained linearly by X

## Linear regression fit



# Estimating the price of a 100 m<sup>2</sup> house in St Kilda

```
new.100 <- data.frame(BuildingArea = 100)
predict(lm.fit, new.100, interval = "confidence")
```


```
##           fit      lwr      upr
## 1 991465.5 894562.7 1088368
```

```
predict(lm.fit, new.100, interval = "prediction")
```

```
##           fit      lwr      upr
## 1 991465.5 13820.26 1969111
```

## Fit improvements

- Remove outliers: black line gives overall fit, blue line fit only to blue data (without red points)



# Linear fit after removing the outliers

```
lm.without.outliers <- lm(Price/1000 ~ BuildingArea, data = st.kilda.data, subset = BuildingArea >= 40)
summary(lm.without.outliers)
```

```
##
## Call:
## lm(formula = Price/1000 ~ BuildingArea, data = st.kilda.data,
##     subset = BuildingArea >= 40)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -876.75 -137.30  -18.27  109.28  896.31
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -289.254     57.471  -5.033 2.22e-06 ***
## BuildingArea   12.305       0.496  24.807 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 298.5 on 97 degrees of freedom
## Multiple R-squared:  0.8638,    Adjusted R-squared:  0.8624
## F-statistic: 615.4 on 1 and 97 DF,  p-value: < 2.2e-16
```

# Extending Simple Linear Regression



- What if I have more than one feature (predictor)?
- House prices depend on more than just BuildingArea! What about
  - land area.
  - Dwelling type (apartment vs unit vs house vs ...)
  - Suburb (location, location, location!)

- Example formula `Response ~ Predictor1 + Predictor2 + Predictor3`
- Left hand side of `~` is the **response** variable (target to predict)
- Right hand side of `~` are the **predictor** variables (features)
- Relationship is assumed to be additive
  - I.e. each additional predictor is added to explain the response  $Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots$
- Interaction or multiplicative terms are denoted with `:` and `*` which are beyond the scope of this course.
  - Would be used to define other relationships
  - E.g.  $Y = \beta_0 + \beta_1 X_1 + \beta_2 X_1 X_2 + \beta_3 X_2 + \dots$

# Multiple linear regression


- Real life problems usually have more than one predictor.
  - Simple linear (single variable) regression can be extended to multiple predictors

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \dots + \beta_p X_p + \varepsilon$$

- The interpretation is the  $\beta_p$  coefficient denotes the average increase/decrease in  $Y$  for each single unit increase in  $X_p$ , holding all the other predictors fixed.

# Extending the house prediction model to multiple features

- Perhaps 100 m<sup>2</sup> houses cost more than 100 m<sup>2</sup> units?



# Multiple regression with `lm`

```
multi.lm <- lm(Price/1000 ~ Type + BuildingArea, data = st.kilda.data)
summary(multi.lm)
```

```
##
## Call:
## lm(formula = Price/1000 ~ Type + BuildingArea, data = st.kilda.data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -700.3  -173.1   -65.9    18.6  3389.6
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   342.865    186.764   1.836  0.06945 .
## Typet        -613.953    286.272  -2.145  0.03448 *
## Typeu        -408.417    139.915  -2.919  0.00436 **
## BuildingArea    9.533     1.014   9.398 2.68e-15 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 469.5 on 97 degrees of freedom
## Multiple R-squared:  0.6989,    Adjusted R-squared:  0.6896
## F-statistic: 75.06 on 3 and 97 DF,  p-value: < 2.2e-16
```

# Model interpretation

```
summary(multi.lm)
```

```
...  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)  
## (Intercept)   342.865    186.764   1.836  0.06945 .  
## Typet        -613.953    286.272  -2.145  0.03448 *  
## Typeu        -408.417    139.915  -2.919  0.00436 **  
## BuildingArea    9.533      1.014   9.398 2.68e-15 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
...
```

```
multi.pred.data <- data.frame(BuildingArea = rep(100, 3), Type = c("u", "t", "h"))  
predict(multi.lm, newdata = multi.pred.data)
```

```
##           1           2           3  
## 887.7252  682.1894 1296.1427
```

# Nonparametric regression or Smoothing

# Parametrics vs non-parametric methods

- **Parametric** methods involve selecting a statistical model (e.g. linear regression model) and fitting the parameters of the model (e.g. slope, intercept) using the training data
- Nonparametric methods don't require selecting a strict model. The data is allowed to *speak for itself*. However, don't have easily interpretable parameters. They are generally intended for **description** rather than formal inference (e.g. k-nearest neighbor smoothing)



# Data smoothing

With **predictor-response** data, the random response variable  $Y$  is assumed to be a **non-linear** function of the predictor variable  $X$ .

$$Y_i = f(X_i) + \varepsilon_i$$

- $f$  is some fixed, non-linear smooth function.
- $\varepsilon_i$  is a zero-mean random variable.
- Smoothing is a **non-parametric method to estimate**  $f$ .

# Local averaging

- Most smoothers (smoothing functions) rely on the concept of *local averaging*
  - In contrast, simple linear regression attempts to fit the best global line.
- E.g. Suppose you want to determine the response  $Y$  conditional on  $x$ .
  - The  $Y_i$  whose corresponding  $x_i$  are near  $x$  should be averaged with higher weight to attempt to estimate  $f(x)$ .
- A generic local-averaging smoother can be written as

$$\hat{f}(x) = \text{average}(Y_i | x_i \in N(x))$$

- where average is some generalised averaging operation.
- $N(x)$  is some neighbourhood of  $x$ .

# Constant-Span Running Mean, k-nearest neighbours

- A simple smoother takes the sample mean of k nearby points
- We define  $N(x_i)$  as  $x_i$  itself, the  $(k-1)/2$  points whose predictor values are nearest below  $x_i$ , and the  $(k-1)/2$  points whose predictor values are nearest above  $x_i$
- This neighbourhood is termed the *symmetric nearest neighbourhood*, and the smoother is called a **moving average** or a **k-nearest neighbours (kNN)** smoother.
- The constant-span running-mean smoother can be written as:

$$\hat{f}(x_i) = \text{mean}[Y_j \text{ such that } \max(i - \frac{k-1}{2}, 1) \leq j \leq \min(i + \frac{k-1}{2}, n)]$$

# Regression splines


- Fit piece-wise functions, where each function can be a  $d$ -dimensional polynomial function
- Constrain the function to be smooth and continuous
- Cubic spline fits cubic polynomial functions, with the constraints:
  - continuous at each knot, continuity of the 1st derivative and continuity of the 2nd derivative
- Advantage of the cubic spline is that the curve looks smooth to the eye, and can be used to fit almost any function

# Loess

- Loess is a Locally weighted scatterplot smoothing method
- The loess (**L**ocal regr**ESS**ion) smoother is a widely used method with good robustness properties.
- It is essentially a weighted running-line smoother, except that each local line is fitted using a robust method rather than least squares.
- As a result, the smoother is nonlinear.
- Loess is computationally intensive and require densely sampled data.

# Local regression


- Fitting local linear fits that are weighted.
- Formula for local constant fit is  $\hat{f}(x) = \frac{\sum_{i=1}^n Y_i K((X_i - x)/h)}{\sum_{i=1}^n K((X_i - x)/h)}$




- Sharp changes in behaviour from the acceleration on a crash test dummy.

# Local regression animation

- Using a large averaging window



- Using a smaller averaging window



# Nonparametric smoothing vs linear regression

- Advantages of non-parametric smoothing
  - Can model non-linear functions (e.g. splines, loess)
  - Does not make any assumption about the functional form of the data
- Advantages of linear regression
  - Computationally efficient, even for multivariate linear regression
  - Model is interpretable, i.e. one can know the statistical meaning of the estimated slope parameters.



## Reference list

James, G., D. Witten, T. Hastie, et al. (2013). *An introduction to statistical learning*. Vol. 112. Springer.