

# Trees and ensemble models

STAT5003

Justin Wishart

## Libraries to load

```
library(tree)
```

## Single tree based methods

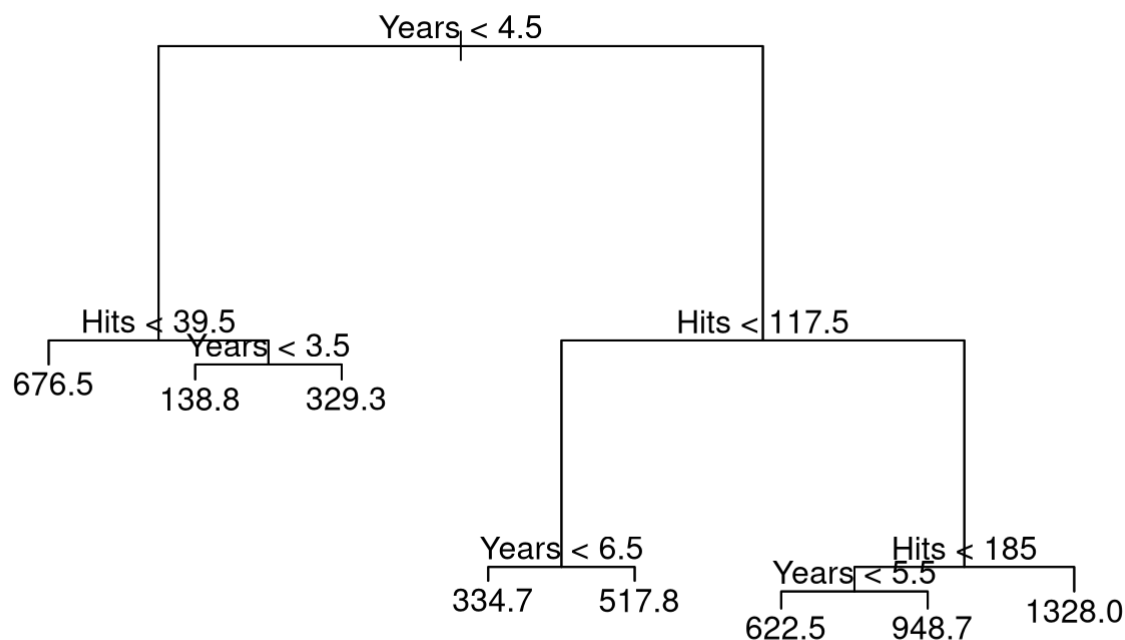
### Regression tree

```
data(Hitters, package = "ISLR")
Hitters <- na.omit(Hitters)
rt <- tree(Salary ~ Hits + Years, data = Hitters)

summary(rt)
```

```
##
## Regression tree:
## tree(formula = Salary ~ Hits + Years, data = Hitters)
## Number of terminal nodes: 8
## Residual mean deviance: 101200 = 25820000 / 255
## Distribution of residuals:
##      Min.   1st Qu.   Median     Mean   3rd Qu.    Max.
## -1238.00  -157.50   -38.84     0.00    76.83   1511.00
```

```
plot(rt)
text(rt)
```



```
# Lecture tree
# rt <- tree(Salary ~ Hits + Years, data = Hitters,
#           control = tree.control(nobs = nrow(Hitters), minsize = 100))
```

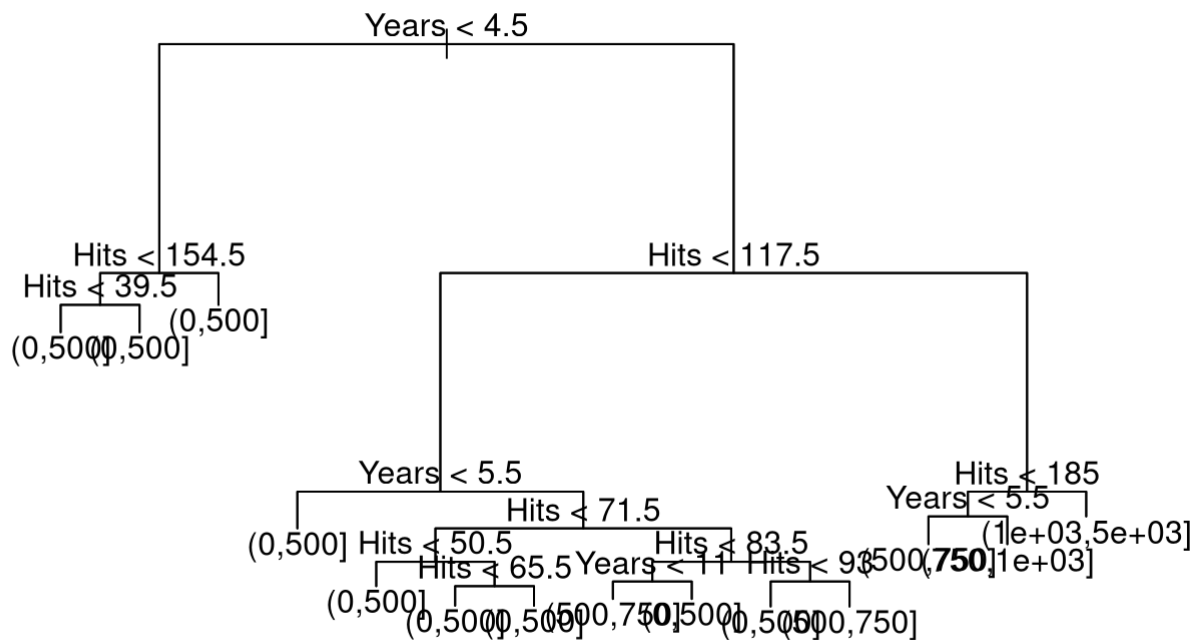
## Classification tree

```
salary.cat <- cut(Hitters[["Salary"]],
                  breaks = c(0, 500, 750, 1000, 5000))
ct <- tree(salary.cat ~ Hits + Years, data = Hitters)

summary(ct)
```

```
##
## Classification tree:
## tree(formula = salary.cat ~ Hits + Years, data = Hitters)
## Number of terminal nodes: 14
## Residual mean deviance: 1.39 = 346 / 249
## Misclassification error rate: 0.308 = 81 / 263
```

```
plot(ct)
text(ct)
```



## Regression tree


This section introduce regression tree using housing value dataset of Boston suburbs

```
library(MASS)
set.seed(1)
train <- sample(1:nrow(Boston), nrow(Boston)/2)

# medv: median value of owner-occupied homes in $1000s.
tree.boston <- tree(medv ~ ., Boston, subset = train)
summary (tree.boston)
```

```
##
## Regression tree:
## tree(formula = medv ~ ., data = Boston, subset = train)
## Variables actually used in tree construction:
## [1] "rm"      "lstat"   "crim"    "age"
## Number of terminal nodes: 7
## Residual mean deviance: 10.38 = 2555 / 246
## Distribution of residuals:
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -10.1800 -1.7770 -0.1775  0.0000  1.9230  16.5800
```


```
plot(tree.boston)
text(tree.boston)
```



```

# check the RSS of the prediction
yhat <- predict(tree.boston, newdata = Boston[-train, ])
boston.test <- Boston[-train, "medv"]
plot(yhat, boston.test)
abline(0, 1)

```



```
mean((yhat - boston.test)^2)
```

```
## [1] 35.28688
```

## Tree ensembles

### Implement bagging ourselves

```
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
set.seed(123)
Hitters[["salary.cat"]] <- salary.cat
Hitters[["binary.salary"]] <- factor(salary.cat == "(0,500]",
                                     labels = c(">500k", "<=500k"))
with(Hitters, table(binary.salary))
```

```
## binary.salary
## >500k <=500k
##    112    151
```

```
with(Hitters, table(salary.cat))
```

```
## salary.cat
##      (0,500]      (500,750]      (750,1e+03] (1e+03,5e+03]
##           151           50           32           30
```

```
inTrain <- createDataPartition(Hitters[["salary.cat"]], p = 0.5)[[1]]
hit.train <- Hitters[inTrain,]
hit.test  <- Hitters[-inTrain,]

# single binary tree classification
tree.model <- tree(binary.salary ~ . - Salary - salary.cat, data = hit.train)
tree.preds <- predict(tree.model, newdata = hit.test)
tree.classified <- levels(Hitters[["binary.salary"]])[apply(tree.preds, 1, which.max)]
tree.accuracy <- mean(tree.classified == hit.test[["binary.salary"]])
tree.accuracy
```

```
## [1] 0.8091603
```

```
# create bagging (a type of ensemble)
n.train <- nrow(hit.train)
bagging.predictions <- vapply(1:100, function(x) {
  idx <- sample(x = 1:n.train, size = n.train, replace = TRUE)
  tree.model <- tree(binary.salary ~ . - Salary - salary.cat,
    data = hit.train[idx, ])
  predict(tree.model, newdata = hit.test)[, "<=500k"]
}, numeric(nrow(hit.test)))

bagging.classified <- ifelse(rowMeans(bagging.predictions) > 0.5, "<=500k", ">500k")
bagging.accuracy <- mean(bagging.classified == hit.test[["binary.salary"]])
bagging.accuracy
```

```
## [1] 0.8549618
```

## Bagging (use ranger package)

```
# Random forests will reduce into bagging if all features are used at every split
# Here we testing bagging by using random forest package and allowing the use of all
# features.
library(ranger)
set.seed(1)

# Bagging for classification
dim(Hitters)
```

```
## [1] 263  22
```

```
names(Hitters)
```

```
## [1] "AtBat"      "Hits"        "HmRun"       "Runs"
## [5] "RBI"         "Walks"       "Years"       "CAtBat"
## [9] "CHits"       "CHmRun"      "CRuns"       "CRBI"
## [13] "CWalks"      "League"      "Division"    "PutOuts"
## [17] "Assists"     "Errors"      "Salary"      "NewLeague"
## [21] "salary.cat"  "binary.salary"
```

```
bag.hit <- ranger(binary.salary ~ . - Salary - salary.cat, data = hit.train,
                  mtry = 19)
summary(bag.hit)
```

```
##               Length Class      Mode
## predictions      132   factor    numeric
## num.trees          1  -none-     numeric
## num.independent.variables 1  -none-     numeric
## mtry               1  -none-     numeric
## min.node.size      1  -none-     numeric
## prediction.error    1  -none-     numeric
## forest             9  ranger.forest list
## confusion.matrix    4   table     numeric
## splitrule           1  -none-     character
## treetype            1  -none-     character
## call               4  -none-     call
## importance.mode     1  -none-     character
## num.samples         1  -none-     numeric
## replace            1  -none-     logical
```

```
mean(predict(bag.hit, data = hit.test)[["predictions"]] == hit.test[["binary.salary"]])
```

```
## [1] 0.8473282
```

```
# Bagging for regression
dim(Boston)
```

```
## [1] 506  14
```

```
idx <- createDataPartition(Boston[["medv"]], p = 0.8)[[1L]]
boston.train <- Boston[idx, ]
boston.test <- Boston[-idx, ]

bag.boston <- ranger(medv ~ ., data = Boston,
                    mtry = 13)
summary(bag.boston)
```

```
##                               Length Class      Mode
## predictions                   506   -none-    numeric
## num.trees                     1     -none-    numeric
## num.independent.variables     1     -none-    numeric
## mtry                          1     -none-    numeric
## min.node.size                 1     -none-    numeric
## prediction.error              1     -none-    numeric
## forest                       7     ranger.forest list
## splitrule                     1     -none-    character
## treetype                     1     -none-    character
## r.squared                     1     -none-    numeric
## call                         4     -none-    call
## importance.mode              1     -none-    character
## num.samples                  1     -none-    numeric
## replace                      1     -none-    logical
```

```
mean((predict(bag.boston, data = boston.test)[["predictions"]] - boston.test)^2)
```

```
## Warning in mean.default((predict(bag.boston, data = boston.test)
## [["predictions"]] - : argument is not numeric or logical: returning NA
```

```
## [1] NA
```

## Random forest

```
set.seed(1)

# Random forest for classification
dim(iris)
```

```
## [1] 150    5
```

```
bag.iris <- ranger(Species ~ ., data = iris, mtry = 1)
print(bag.iris)
```

```
## Ranger result
##
## Call:
##  ranger(Species ~ ., data = iris, mtry = 1)
##
## Type:                               Classification
## Number of trees:                     500
## Sample size:                         150
## Number of independent variables:     4
## Mtry:                                1
## Target node size:                    1
## Variable importance mode:            none
## Splitrule:                           gini
## OOB prediction error:                 4.67 %
```



```
# Random forest for regression  
dim(Boston)
```

```
## [1] 506 14
```

```
rf.boston <- ranger(medv~., data = Boston, subset = train, mtry = 6)
```


```
## Warning in ranger(medv ~ ., data = Boston, subset = train, mtry = 6): Unused  
## arguments: subset
```

## Boosting

```
# regression=  
library(gbm)
```


```
## Loaded gbm 2.1.8.1
```

```
set.seed(1)  
boost.boston <- gbm(medv ~ ., data = Boston[train, ],  
                    distribution = "gaussian", n.trees = 5000)  
summary(boost.boston)
```



```
##          var      rel.inf
## rm          rm 34.5235940
## lstat      lstat 32.0730829
## dis        dis  7.9394616
## crim       crim  5.5339488
## black      black 4.0130690
## ptratio    ptratio 3.7836978
## age        age  3.7646124
## nox        nox  3.6395377
## tax        tax  1.8703031
## indus      indus 1.2439324
## rad        rad  0.9199469
## chas       chas  0.4607448
## zn         zn   0.2340686
```

```
# classification
adaBoost.model <- gbm(as.numeric(binary.salary) ~ 1 ~ . - Salary - salary.cat,
                      data = hit.train, distribution = "adaboost",
                      n.trees = 5000)
summary(adaBoost.model)
```



```
##           var      rel.inf
## CHits      CHits 21.56041812
## Hits       Hits 15.51432443
## CRBI       CRBI 11.30231278
## CRuns      CRuns 11.24978881
## CWalks     CWalks 8.88705767
## PutOuts    PutOuts 7.57615788
## Years      Years 4.11295408
## AtBat      AtBat 3.16141049
## Runs       Runs 2.90773123
## Assists    Assists 2.39679154
## HmRun      HmRun 2.38027999
## NewLeague  NewLeague 2.22759001
## RBI        RBI 1.78945279
## CAtBat     CAtBat 1.38958898
## Walks      Walks 1.06865094
## CHmRun     CHmRun 1.00563440
## Errors     Errors 0.74739700
## Division   Division 0.68238686
## League     League 0.04007202
```

```
preds <- predict(adaBoost.model, newdata = hit.test,
                  n.trees = 5000, type = "response")
mean(ifelse(preds > 0.5, "<=500k", ">500k") == hit.test[["binary.salary"]])
```

```
## [1] 0.8167939
```